

Swift Racks Product Innovation Guide — MVP Edition (Parts 0–3)

Purpose

- A lightweight playbook for PMs, founders, and beginner and higher devs building AI-powered B2B SaaS.
- Takes you from deciding what to build → building a thin-vertical MVP fast → launching a private beta → learning weekly.
- Designed to be used with Swift CNS (our AI-powered innovation platform). You can revise labels later to fit your processes.

How to use this document

- Read Part 0 once. Then work through Part 1 → Part 2 → Part 3 in order.
- Each section lists: Purpose, Outcome, Inputs, Activities (outline), Artifacts, Swift CNS, Exit Criteria.
- Default cadence: Part 1 = 2–3 days, Part 2 = 3–5 days, Part 3 setup = 1–2 days, then weekly learning loop.

Core artifacts for the MVP (names only)

- Opportunity/Idea Brief, Assumption Table, Hypothesis Card, Experiment Card, Learning Log, Decision Record, MVP Slice Map, Prompt/Eval Sheet, Release Checklist, Pilot Success Plan, Tracking Spec (lite), Changelog.

PART 0 — START HERE

Welcome & purpose

- Set expectations and show the path: Decide → Build → Launch. You will ship a problem-validated MVP to a small, high-signal cohort.

Who this is for

- Product managers, founders, beginner developers focused on AI B2B SaaS.

What you'll ship

- A thin-vertical MVP instrumented with a North Star Metric (NSM) and 2 input metrics, plus a private beta plan.

How to work with Swift CNS (at a glance)

- One Discovery board (Idea → Assumptions → Hypotheses → Experiments → Learnings → Decision).
- One Build board (MVP Slice → Build Tasks → QA/Checks → Ready).
- One Launch board (Pilots → Experiments → Decisions → Changelog).
- Use the mini-templates listed under “Core artifacts.”

Access and governance

- Manual access for private readers. Weekly changelog note. Gate A (Build Go/No-Go) and Gate B (Launch Ready).

Cadence and gates

- Gate A after Part 1 (evidence + MVP slice defined).
- Gate B after Part 2 (quality checks + telemetry + rollback).

Program paths

- PM track, Dev track, Fast-track (2–4 week MVP).

PART 1 — DECIDE WHAT TO BUILD (SWIFT CNS INNOVATION LOOP)

Overview

- Goal: converge on one crisp problem, three critical assumptions, 1–2 testable hypotheses, one smallest possible experiment, and a clear decision.

Steps at a glance

1. Idea / Real-World Problem
2. Identify Critical Assumptions
3. Form Testable Hypotheses
4. Design & Run Experiments
5. Extract Key Learnings
6. Synthesize Insights → Decision (Proceed, Pivot, Stop)

/decide/01 — Idea / Real-World Problem

- Purpose: define a sharp problem for a specific ICP.
- Outcome: one-sentence problem statement supported by at least one real signal.
- Inputs: ICP guess, quotes, tickets, emails, screenshots.
- Activities: capture context and pain; tighten to one sentence.

- Artifacts: Idea/Opportunity Brief.
- Swift CNS: create an Idea card with fields for ICP, problem, evidence.
- Exit criteria: you can read the problem in one breath; at least one real-world quote/signal attached.

/decide/02 — Identify Critical Assumptions

- Purpose: surface and rank the riskiest assumptions.
- Outcome: top three assumptions with a proposed test each.
- Inputs: Idea Brief.
- Activities: list assumptions across desirability, feasibility, viability, usability/safety; rank by Impact × Uncertainty.
- Artifacts: Assumption Table.
- Swift CNS: add assumptions to the Idea card; sort by risk.
- Exit criteria: exactly three prioritized assumptions, each with a test idea and timebox.

/decide/03 — Form Testable Hypotheses

- Purpose: turn assumptions into falsifiable bets.
- Outcome: 1–2 hypotheses with metric, threshold, and time window.
- Inputs: Assumption Table.
- Activities: write hypotheses; define kill criteria.
- Artifacts: Hypothesis Card(s).
- Swift CNS: create Hypothesis cards linked to the Idea.
- Exit criteria: each hypothesis has metric + threshold + time; stop rule written down.

/decide/04 — Design & Run Experiments

- Purpose: run the smallest test that can fail.
- Outcome: one experiment per active hypothesis.
- Inputs: Hypothesis Cards.
- Activities: pick method (interviews-5, landing page + CTA, concierge/manual, prototype usability, data pull); define success and guardrails; execute.
- Artifacts: Experiment Card with plan, sample, and evidence.
- Swift CNS: create Experiment cards; move status Planned → Running → Completed.
- Exit criteria: evidence attached and result marked Pass/Fail/Inconclusive.

/decide/05 — Extract Key Learnings

- Purpose: convert evidence into facts and patterns.

- Outcome: 3+ high-confidence learnings tied to metrics.
- Inputs: experiment evidence.
- Activities: capture observations; separate facts from interpretations; tag themes; note impact on NSM/input metrics.
- Artifacts: Learning Log.
- Swift CNS: add learnings with confidence and links back to evidence.
- Exit criteria: at least three learnings with confidence and metric impact.

/decide/06 — Synthesize Insights → Decision

- Purpose: make a clear call and define the MVP slice.
 - Outcome: Decision = Proceed, Pivot, or Stop; if Proceed, an MVP Slice Map and initial metrics are defined.
 - Inputs: Hypothesis Cards, Experiment results, Learning Log.
 - Activities: summarize rationale; note risks and mitigations; define the golden path for MVP.
 - Artifacts: Decision Record, MVP Slice Map (value moment, steps, NSM + two inputs).
 - Swift CNS: record Decision; if Proceed, auto-create MVP card on the Build board.
 - Gate A (Build Go/No-Go): proceed only with strong evidence (e.g., hair-on-fire quote or 10 qualified signups or paid LOI), a passing hypothesis or a justified pivot, and a defined MVP Slice + NSM.
-

PART 2 — BUILD THE MVP FAST (THIN VERTICAL)

Overview

- Goal: stand up a single golden path delivering value end-to-end, with basic quality, telemetry, and rollback.

Pages (05–08)

- 05 Setup
- 06 AI Baseline
- 07 UX Thin Vertical
- 08 Quality & Observability (Lite)

/build/05 — Setup

- Purpose: create a working skeleton path and developer basics.
- Outcome: repo ready, simple environment, org login/basic auth, seed data, feature flag.
- Inputs: MVP Slice Map.

- Activities: scaffold repo; implement non-AI golden path; enable feature flag; seed demo data.
- Artifacts: Coding Charter (one page), Release Checklist (MVP).
- Swift CNS: Build board ticket(s) for setup tasks; link to checklist.
- Exit criteria: you can demo the golden path without AI; checklist started.

/build/06 — AI Baseline

- Purpose: connect the smallest useful AI capability to the golden path.
- Outcome: one realistic input produces one useful output; basic evaluation recorded.
- Inputs: golden path skeleton; sample inputs/outputs.
- Activities: choose RAG or prompt-only; wire evaluation example; log prompts, latency, and cost.
- Artifacts: Prompt/Eval Sheet (tiny).
- Swift CNS: track prompts/evals; link costs/latency to the Build card.
- Exit criteria: deterministic demo works; simple eval captured.

/build/07 — UX Thin Vertical

- Purpose: ship minimal UI that delivers the value moment.
- Outcome: one page or two screens implementing the golden path with clear empty/error states.
- Inputs: MVP Slice Map.
- Activities: wireframe → implement; write concise copy; add loading/retry patterns.
- Artifacts: Golden Path Wireframe.
- Swift CNS: link design to Build tickets; connect to feedback intake.
- Exit criteria: a new user can complete the value moment on first attempt.

/build/08 — Quality & Observability (Lite)

- Purpose: basic safety net and visibility before beta.
 - Outcome: smoke tests, feature flag/kill switch, event logging for NSM + inputs, cost per action logging.
 - Inputs: working golden path.
 - Activities: add minimal tests; instrument key events; write rollback steps.
 - Artifacts: Test & Telemetry Mini-Spec, Release Checklist (completed).
 - Swift CNS: events visible in Launch dashboards; attach runbook.
 - Gate B (Launch Ready): rollback plan exists; logs verify NSM events; smoke tests pass.
-

PART 3 — LAUNCH, LEARN, ITERATE (PRIVATE BETA)

Overview

- Goal: get 3–10 qualified accounts using the MVP, see real behavior, and improve weekly.

/launch/09 — Private Beta

- Purpose: recruit a small, high-signal cohort and get them live.
- Outcome: at least three ICP accounts live.
- Inputs: ICP list, MVP build.
- Activities: select accounts; define success plan per account; manual access; onboarding checklist; support channel.
- Artifacts: Pilot Success Plan (per account), invite email snippet, onboarding checklist.
- Swift CNS: Pilot tracker with per-account goals and status.
- Exit criteria: three accounts live this week.

/launch/10 — Instrumentation

- Purpose: see what happened yesterday in two minutes.
- Outcome: dashboards for NSM, two input metrics, daily active pilots, and cost per value event.
- Inputs: event schema from Part 2.
- Activities: finalize a short tracking spec; wire simple dashboards.
- Artifacts: Tracking Spec (lite).
- Swift CNS: metrics dashboard for pilots.
- Exit criteria: yesterday's numbers are visible and trustworthy.

/launch/11 — Experiments

- Purpose: run one learning loop per week.
- Outcome: one experiment shipped weekly, tied to a metric.
- Inputs: pilot feedback and usage data.
- Activities: create an Experiment Card; run; log decision; add to changelog.
- Artifacts: Experiment Card, Decision Log entry.
- Swift CNS: Experiment board with weekly cadence.
- Exit criteria: one metric-linked change shipped each week.

/launch/12 — Feedback Ops

- Purpose: capture VOC without chaos and turn it into decisions.
- Outcome: every thread is tagged and linked to a decision or backlog item.

- Inputs: pilot conversations, support tickets, interviews.
- Activities: set up one intake form + tags; weekly triage; monthly NPS/PMF pulse (very short).
- Artifacts: Feedback Triage Board, Changelog template.
- Swift CNS: feedback tags and rollups connected to experiments and roadmap.
- Exit criteria: zero orphaned feedback; changelog updated weekly.

WHAT'S DELIBERATELY OUT OF SCOPE FOR THIS MVP

- Deep pricing and packaging work, advanced architecture/performance, SOC2 and enterprise questionnaires, complex analytics taxonomies, multi-region resilience, large template libraries, full Swift CNS how-to manuals. Add these after beta feedback.
-

DEFAULT ONE-WEEK PRODUCTION PLAN (FOR PUBLISHING THE GUIDE)

- Day 1: Write Part 0. Create blank templates for the core artifacts.
- Day 2: Write Part 1 sections (01–06) at outline depth. Add Idea, Assumption, Hypothesis, Experiment templates to CNS.
- Day 3: Write Part 2 sections (05–08). Add Release Checklist and Prompt/Eval Sheet to your repo template.
- Day 4: Write Part 3 sections (09–12). Create feedback intake and a simple dashboard for NSM + inputs.
- Day 5: Add one tiny example to each page, edit lightly, publish privately, enable manual access.
- Weekend or next Monday: Run one real idea through the loop; tighten anything confusing.